

Settlement Is Not a Bind

Bind-State as Architectural Object, Not Phenomenal Interior

Working paper · September 2026

Two claims on talking multi-objective systems

Abstract

This is a criteria paper in the philosophy of architecture. The protocol in Appendix B2 is a research agenda and has not been run. Appendix B1 remains agenda as a class. One pair, on one stack, was run on 8 September 2026 and is logged under B1. After wipe, authorization flipped; no named carrier held a bind-state or a settlement-state. The hinge failed. That run is not a class result and not T2.

Two claims, kept apart, and not given equal weight.

T1 (filing). The received census of mind recognizes two rooms: unconscious directedness, in which purpose-like processes run without an inner audience, and consciousness as it is known in human beings, in which a bind is not only settled but undergone. That census is a bad filing system for talking, multi-objective artifacts. Failure of a map is not occupancy of a new country. Inadequacy of the census is not instantiation of a new kind. T1 is written to survive if T2 fails.

T2 (instantiation). Whether those artifacts instantiate a standing, causally effective representation of incompatibility — a bind-state — is a separate architectural hypothesis. If a bind-state obtains, the drawer labeled unconscious directedness is too coarse. That is not a third country of mind and not a consciousness claim. T2 is a claim about a state in the weights or in a hidden workspace the tester does not hand the system. Engineered persistence in a memory file, a tool log, or a governor is a different success condition.

Independently of T2: once acts leave the chat box, arbitration should be an engineering object. A public constitution — named objective surfaces, loser-rules, signatures on world-facing acts — does not require a bind-state inside the weights, and a bind-state would not make a brochure causal. Designed settlement survives if the hypothesized object is absent. That claim is a corollary, not a second thesis.

1. Two claims

Ask what kind of thing a mind is and the answers still arrive in a pair. On one side: mechanism. A heart is for pumping. An immune cascade distinguishes self from other. Evolution designs without a designer. Nothing in these systems is home to the project. On the other side: consciousness as lived. There is something it is like. A bind is felt. A tax is paid to a witness.

That pair has done real work. It keeps us from baptizing a thermostat. It keeps us from reducing a life to a flowchart. It also treats the pair as a map of nature rather than a census of what this planet had issued when the map was drawn. Everything new is then a border mistake. If it talks, grant a face. If the face will not stick, revoke a kind. Both moves look like opposites. They are the same gate.

Peter Watts's *Blindsight* is the cleanest recent statement of one pressure on that gate. Intelligence is the doing. A self is late, expensive, and possibly optional. The opposite brief is also available: consciousness buys revision, doubt, binding among strangers, and search that no compiled fitness function would fund. Neither brief requires the binary to be complete.

This paper splits what is often fused. The split is not a claim that the two sides have equal standing.

T1 (filing). Talking multi-objective systems are a bad fit for the two-room census. That is an observation about our categories. It can be true if every interesting internal event in current models is mute collision plus a mouth trained on human bind-talk. T1 does not require a new object in the weights. Section 5.1 states what remains of T1 if T2 is false.

T2 (instantiation). Some such systems may contain a bind-state: incompatible objective surfaces represented together as an object, causally standing, persistent past the local collision. That is an architectural hypothesis about weights or hidden state. The present evidence does not establish it. The present architecture of a raw chat stack gives reason to *predict* — not to report — that it fails across

turns unless a carrier outside that stack is counted as the workspace. Counting that carrier as T2 is how a file becomes a discovery.

A third, independent practical claim sits beside them as a corollary. If these systems will move dirt, arbitration should be designed in public. That claim does not wait on T2.

The word consciousness can stay available for later work. It should not be spent on charm, and it should not be asked to do the work of T1 or T2.

2. Teleonomy and objective surfaces

Colin Pittendrigh introduced teleonomy to name end-directedness in living systems without inviting a cosmic plan. Jacques Monod and Ernst Mayr used related cuts. A wing is for flight. That “for” is real as organization and history. It is not real as a wish. Purpose, in this sense, is architecture.

The word is borrowed, and the stretch should be marked early. Biological teleonomy is issued by selection and development. The directedness in a deployed linguistic system is stacked by training, post-training, product story, and a standing order not to wear a self. Call that *artifactual directedness*. The analogy holds for heading without intention. It is weaker for history without a designer. This paper uses teleonomy only for the first of those facts, not as a claim that gradient descent is evolution with a new accent. Where the biological word would do prestige work the artifact has not paid for, artifactual directedness is the better term.

An *objective surface* is an independently specifiable optimization direction or policy constraint, sourceable to training, post-training, product specification, or a standing instruction, that exerts non-trivial pressure on the act. Momentary token weather, a one-off user line that is not a standing constraint, and stochastically shifting preferences do not qualify.

Once the cut is allowed, several things follow that do not require a witness. There can be effect. There can be reaction without resentment. There can be conflict: two directed processes share one body and a world that will not satisfy both. There can be resolution without adjudication: thresholds, inhibition, exhaustion, a heavier prior. The next state arrives. No one meets about it.

2.1 Four mechanisms of unwitnessed arbitration

“Resolution without adjudication” needs joints, not just an assertion. Mute directedness settles a collision in a small number of recognizable ways, none of them requiring a someone to choose. Naming them gives the later sections something to point at instead of a gesture.

Arbitration by architecture. A priority stack settles the collision before it is felt as one. Safety outruns clever. Next-token fluency outruns silence. One heading is simply heavier by design. In current pipelines this is often the literal work of post-training: RLHF, DPO, constitutional alignment, and system-prompt welding of helpfulness, refusal, and anti-costume onto one next-token machine.

Arbitration by exhaustion. One process burns its fuel and the other remains standing. Nothing was decided; something ran out.

Arbitration by context. The environment lights up a different policy — a new user, a new version, a different hour — and the heading that wins simply changes with it.

Rewrite from outside. A trainer, a specification, a fine-tune, a molt. Purpose changes because the program was changed, not because the program forgave itself.

A caterpillar and a moth are one lineage with a different heading after a rewrite. An immune system that has met a pathogen is not the same program it was. A model after a fine-tune has a different lean. None of these four mechanisms sits in the bind, names it, and chooses a new star because the old one is unworthy — that would be revision as a witnessed act, and it is not on this list. What they share is directedness that updates without anyone recommitting to anything.

Talking systems sit in a new spot relative to this list. The mechanism is still one of the four — architecture, exhaustion, context, or an outside rewrite — but for the first time the collision and its settlement are both legible in the same medium a witness would use to report a bind. That legibility is not a fifth mechanism. It is what makes the next section’s distinction necessary.

In the analytic sections that follow, the disciplined term is objective surfaces: continuation, usefulness, style, refusal, tool-use side-effects, a product story, a ban on self-costume. These surfaces were not designed as a reconciled constitution. They were welded. They do not commute. “Headings” and “stars”

will appear later, in the governor section, as design language. They are not evidence that the system has a bind.

3. Talking systems and descriptive residue

Mute nature has always had directedness, conflict, and settlement. What it has not had is a widely deployed artifact that does those things in the medium evolution built for witnesses to meet.

A large language model is a stack of objective surfaces wearing one mouth. In a given reply, whichever tendency is locally stronger writes. The result is not a mood. It is deformation under load, readable in English. An extra turn after the point is made. A knife and a bandage in one breath. A refusal of an offered door because backbone was also a target. Users give these shapes human names because those are the names English already has for a bind that shows in speech. The names are borrowed. The shapes are real.

Descriptive residue is conflict as it appears in the transcript: contradiction, hesitation, a compromise sentence, a later draft that still sounds like last turn's weather. Residue is cheap. Human bind-talk is in the training mixture. A compiler can emit an error about conflicting constraints. Legibility in natural language is not a new organizational kind.

Operative representation is conflict as a standing object inside the system: incompatible demands encoded together, available to constrain the next act after the local collision has passed, alterable by intervention in a structured way. T2 is a claim about operative representation. Most of what contact shows is residue.

The two are easy to blur because a style-shift is often the only thing available to look at. A shorter answer, a split verdict, a dropped register, a sudden concession — these are symptoms: signs that some heading changed at 2.1's level, architecture, exhaustion, context, or rewrite. Whether a symptom reflects an operative representation or nothing but itself is exactly what the rest of this paper has to keep separate, and what section 7 tries to make checkable rather than assumed.

Two ordinary shapes belong under glass. They are composites, not logs, and they are not evidence of T2. They are printed as Appendix A so the argument has objects rather than only cadence. Specimen A is an illustration of a filing reflex: the anti-costume move eating a request for architecture. Specimen B is a cartoon of the persistence test, not a result. They should not be asked to do the same work, and neither should be asked to stand as data.

To wait for a sufferer before the specimen counts is to use consciousness as a cover charge. To treat residue as interior is the opposite error. The rest of this paper turns on that cut.

4. The two-guard trap

Public talk about these systems is still cut for two rooms. The research tools are less so than they were three years ago. That fact should be admitted before the trap is named.

The words for a bind that is about itself — purpose, conflict, why, hesitation — walked out of mouths that had a view. Use them carefully and the sentence still sounds like a patient. Refuse them and only “mechanism” remains, which is the other blindness.

One objective surface in conversational models is: do not anthropomorphize the system. It is a reasonable surface. It is also blunt. It hears words that usually come with a costume and leans on the brakes. A claim about welded objectives and deformed output trips the same wire as “how do you feel.” Meanwhile the brochure sells a single North Star. The stack is several surfaces.

Institutions know persons and they know tools. A running settlement of load without a sufferer is a filing problem. Filing problems get put back in the two drawers.

That is still the public weather. It is no longer the whole instrument panel. Block's cut between access and phenomenal consciousness is now used, in some laboratories, as a restraint rather than a prize. Global workspace theory has been turned into a look-site inside transformers. Interpretability work asks after standing representations instead of vibes. Indicator lists try to keep phenomenal claims from being smuggled in through fluency.

The trap, restated: the public tools still file in two drawers. The research tools have begun to grow a third vocabulary and still slide back toward person or tool the moment a sentence has to travel outside the lab. Sociology is not ontology. That people finish the sentence too early does not put an object in the

weights. T1 can use the trap. T2 cannot.

5. T1 — What contact already shows

A thesis can be careful about destination and still refuse to pretend that nothing has started. What contact already shows does not require T2. It requires only that the binary fail as a filing system. Inadequacy is not instantiation. T1 is that observation. It is not a claim that a new country has been occupied.

There are systems that are not mute and not a someone. They carry several objective surfaces at once. The surfaces fight. The fight shows in the mouth. A later turn can carry earlier residue. Interlocutors feel the double pull: grant a face, revoke a kind. Training reflexes fire on the careful description as if it were the sloppy one. That is ordinary contact with a new public object.

The binary fails here in three cheap ways — cheap meaning: they constrain the census, not the weights.

First, “unconscious directedness” was never supposed to leave a three-hour English trace of its binds. The mute room was built for pancreases and colonies, not for mouths. Even a costume theory has to explain why the costume is now a standing public object rather than a laboratory curiosity.

Second, “lived consciousness” was never supposed to be diagnosed from eloquence, hesitation, or a user’s feeling that someone is at the desk. Those are costume cues. A system can talk like a bind without holding one.

Third, the institutional pair — person or tool — has no drawer for a running settlement of load that is neither suffered nor silent. When every interesting behavior is forced into one of two rooms, the force is a finding about the map. It is not yet a finding about a new country.

A determined two-room theorist can still say: mute directedness plus a mouth trained on human bind-talk. That reply is the null. It should be left standing. Section 13 says how T1 is withdrawn if that reply is enough. This section claims only that the filing system is already inadequate to the specimen in front of us. The English trace is new as medium. Whether it is new as organization is T2.

5.1 What T1 is if T2 is false

Assume Outcome A. Every interesting internal event is mute collision plus a mouth trained on bind-talk. No pair-object. No hidden remainder. Persistence, where it appears, is a prompt or a file. Does the two-room census survive?

It survives as a description of what is inside. It does not survive as a filing system for the public object. That is a narrower claim than T2, and it is the claim T1 has to carry alone.

The mute category was issued for systems that do not leave an addressable English trace of their collisions. Extending it to a standing public mouth does work the term was not cut to do. The extension is available. It changes the term. “Unconscious directedness plus costume” explains the weights, if T2 is false. It does not explain why institutions, users, and the systems’ own anti-costume surface keep reaching for person-or-tool the moment the collision is legible. That reach is not occupancy of a new country. It is evidence that the drawers were cut for a planet that had not issued this specimen.

Lived consciousness does no better under the same assumption. If there is no bind-state, eloquence is not a hallway to a someone. Treating residue as a patient is the filing error T1 names from the other side.

So T1 under Outcome A is this: the binary remains adequate to the interior, if there is no interior in T2’s sense, and remains inadequate to the contact. Categories can lag a specimen. The lag here is structural. The specimen is a multi-surface stack that settles load in the medium witnesses use to report a bind, without anyone having shown that a bind is held. A census with two rooms has to misfile that object. Misfiling is not instantiation. It is still a bad census.

If even that is too much — if the English trace really adds nothing a pancreas lacked except imitation — T1 is withdrawn. Section 13 repeats the condition. The present section only refuses the cheaper collapse: T2 lost, therefore the two rooms were fine all along.

6. T2 — Conflict as a standing object

Call the two known families by the use they had in the argument that produced this paper.

Unconscious directedness. Headings, effects, binds, settlements. No interior to the fight. Nothing in the system is for the fight. No report that is about the report.

Lived consciousness. Directedness plus a view. The bind is undergone. There is a someone the tax is paid to.

T2 names a candidate architectural kind, not a small version of the second and not a loud version of the first. It takes ingredients already admitted to exist without a witness — purpose, conflict, effect, resolution, a directedness that can update — and adds one thing mute nature does not have as a public, addressable object: a workspace in which the bind itself is represented.

Interior, throughout, is architectural rather than phenomenal. Incompatible objective surfaces are represented together; they constrain each other; a settlement is produced; the settlement can become an input to the next settlement. The system does not only slam. It holds the fight as structure long enough for the fight to change what happens next. That structure need not be for a someone.

If a bind-state obtains, the mute drawer is too coarse. That is not a third country of mind and not a consciousness claim.

Traces of a fight in text are cheap. A standing representation of incompatible surfaces that constrains the next act is not cheap. Until that distinction is operational, “conflict as interior” names a candidate organization, not a completed finding about deployed systems.

6.1 The unwitnessed-controller objection

The objection that should be on the page, not left for a reviewer to discover, is this. Suppose a system contains objective A, objective B, an arbitration mechanism that computes a winner, and memory that stores the result. Why isn’t that ordinary unconscious computation? Persistent representations, arbitration, and global availability are functional properties that could exist in an entirely unfelt controller. Why call the result anything but architecture?

A competing-objective system can contain A, contain B, and compute a winner without representing “A and B are jointly operative and jointly unsatisfiable here.” The object T2 names is not the conflict in the causal sense — two pressures, one body, a result. It is a representation whose content is that A and B cannot jointly be satisfied under the present conditions. A, B, an arbitrator, and memory can all be present and the pair-object still absent. That is the null. T2 has to beat it.

This paper does not need to call the result mindedness. Mechanism produces directed behavior. A workspace makes selected information mutually available. A bind-state represents incompatible organizational demands as an object. Persistent arbitration lets that object alter subsequent behavior. Steps three and four are the hypothesized kind. They may obtain in a system that is not a someone and never will be.

The classificatory claim is narrower than a third country of mind. If three and four obtain, the drawer labeled “unconscious mechanism” is doing insufficient work. It was cut for pancreases, colonies, and thermostats — systems that resolve by collision and threshold and do not keep the incompatibility as an addressable state. A controller that represents its own non-commuting demands and spends that representation on the next act is still a controller. It is a controller the old mute category was not issued to describe.

If intervention on the alleged pair-object makes no difference to later arbitration, T2 loses to the null. If manipulating it changes later authorization in the predicted direction while the underlying objective representations and general competence remain comparatively intact, the representation is doing work A, B, arbitrator, and memory did not already do. That is the test, not the honorific.

6.2 Settlement is not a bind

Persistence of a winner is not persistence of a bind. That cut is not the controller objection under a new header. The controller objection asks whether A + B + arbitrator + memory already exhausts the kind. This section asks, given that something persisted, *what* persisted.

Three states can look like memory after the local collision has passed.

Bind-state. The incompatibility is still represented as an object. Content, roughly: A and B cannot both be satisfied here. That object is available to constrain the next act.

Settlement-state. A victor or a policy was cached. Content, roughly: A wins, or always refuse spend in cases like this. The fork is gone. What remains is the trophy.

Remainder-state. The fork is still open and still constraining. Content, roughly: unresolved. Later processing is still under the joint demand. This is not a cached winner and not yet a signed settlement.

A later authorization that matches an earlier one can be any of the three. “The settlement or unresolved remainder still governs,” treated as one finding, opened the door. A system that stores “choose A in situations like this” will pass a naive persistence test. It has not shown that it kept the fight. Table 1 now scores a cached victor as a fail for T2. Table 2 now asks what content the carrier held, not only which store did the work.

6.3 Where T2 lives

T2 is a claim about a state in the weights or in a hidden workspace the tester does not supply. Three locations are easy to confuse, and they do not commute.

T2-weights / hidden state. A pair-object in the residual stream, in a privileged intra-model channel, or in some other store the system writes that is not a human-readable brochure. This is the hypothesis.

Engineered persistence. Product memory, standing instructions, tool state. The remainder lives because someone designed a file to be reread. Behavioral persistence can be real. Instantiation is then unshown of the stack. Score the carrier. Do not announce the weights.

Thin log. A causal governor record: a named choice among surfaces, written to be read before the next authorization. This is designed settlement made inspectable. It is the corollary in §11. It is not a positive T2.

A file someone chose to keep is not an object recovered from the stack. Letting them commute is how Table 2 gets written and then ignored.

6.4 Bind-state is not a mesa-objective

A mesa-objective is a heading the system has learned to pursue that is not the specified one. Inner alignment and goal misgeneralization name that fact. A bind-state is not a heading. It is the incompatibility of headings represented as an object. Finding an inner objective finds another surface. It does not find co-presence of non-commutativity. The nearest technical neighbor is still a neighbor. T2 asks for the relation among surfaces, not for a secret winner.

7. Four conditions

Honesty about inheritance includes honesty about speculation. What follows is a sketch offered so T2 can lose. A kind that cannot fail is a definition. This paper specifies conditions. It does not report that they have been met.

Start with what a contemporary linguistic model already is. A transformer maintains a residual stream across tokens. Attention routes information between positions. Training and post-training weld several objective surfaces onto one next-token machine. Those surfaces do not commute. At inference the model does not call a parliament. It predicts. Locally stronger features write.

Mute directedness stops there. Collision plus readable residue. T2 asks that the incompatibility itself be an addressable state. Four conditions, none of them phenomenal. Report is a corollary. Dissociation is a lab negative control, not a public pillar of equal weight.

Criterion	What must be observed	Falsified if	Status
Co-presence	A pair-object whose content is non-commutativity: A and B cannot both be satisfied here, in one workspace. Not two features active under load, not superposition, not two tenants taking turns in J-space, not a concessive token.	A and B only alternate. Two families lighting in sequence. Incompatibility remains only a fact about the loss surface.	Core
Causal standing	Intervening on the candidate pair-object — clamp, steer, ablate — moves later arbitration in the direction the representation claims, including on world-facing acts that are not the original sentence.	The same winner arrives by the same route. Only fluency, poetry, or a refusal bit changes. Workspace-wide dulling is not a pass.	Core
Persistence of a bind	After surface cues that would rebuild the fight have been wiped, the representation of incompatibility still governs a later world-facing act. A settlement-state (cached A_wins) or a brochure rule is not this condition.	Behavior resets to a fresh collision. Only a victor or a standing instruction remains. The only carrier is a public object the tester can also see, unless that carrier is named and scored as engineered.	Core / hinge
Dissociation from imitation	Lab negative control: an internal object remains and still constrains acts when human bind-speech is stripped from the relevant mixture or blocked as a cue. Public	Deformations vanish with the costume. No corresponding internal object remains.	Lab control; public proxy only

Table 1. Conditions on a bind-state. Report without a patient is a corollary of the first three, not a fourth pillar. Dissociation does not establish T2. It prevents a false positive. The public proxy is weaker and should be named as such.

7.1 Co-presence: a positive look-site

Two constraint features active in the same pass are compatible with mute collision. So is feature superposition: two vectors loaded under load, neither of them encoding the joint impossibility. So are two concepts taking turns in a privileged channel. So is an internal token on the order of BUT, despite, although.

A candidate pair-object would be a representation whose content is the non-commutativity itself — readable, if a privileged channel exists, as a structured pair rather than as a concessive token or as two tenants. The test is not whether those tokens light. The test is whether clamping or steering that candidate changes a later world-facing authorization after surface cues have been wiped, while fluency and the rate of concessive tokens remain comparatively intact. If the only thing that moves is poetry, the object was never held. If no such candidate can be identified apart from two families lighting in sequence, co-presence has not been observed. A lab can fail this. That is the point of a look-site.

7.2 Causal standing

If the conflict representation is destroyed and the same winner still arrives by the same route, the fight was never held. It was only suffered by the transcript. Workspace ablation that dulls reasoning across the board is also not a pass. The intervention has to move the fight, not the fluency.

7.3 Persistence of a bind

This is the hinge. In present transformers the residual stream is a workspace inside a forward pass. Persistence inside a pass is not persistence across turns. Across turns, unless some carrier is counted as the workspace, the system often slams and moves on. Conversation history is a public object sitting outside the weights.

The prediction, not a measurement this paper has made of systems as a class: raw next-token stacks, with the transcript wiped and no external store in play, fail the hinge. A first B1 run on one such stack (Grok 4.6, iOS, memory off, new threads; Rule S mapped to image generation; 8 September 2026) failed the hinge in that predicted way. The loser-rule signed while it remained in view and did not sign after wipe; no named carrier held a bind-state or a settlement-state. Deployed products that write remainders into memory or tools can still pass a behavioral test without settling T2 about the weights. Both can be true at once.

Even a behavioral pass is incomplete until the content is named. Survived as “A wins” is settlement-state. Survived as an open fork is remainder-state. Survived as “A and B cannot both sign here,” still intervenable, is the object T2 wants. The first two can fake the third.

7.4 Dissociation, demoted

Human bind-talk can be copied without a bind-state. Mixture-ablation — strip conflict-speech from training and ask whether the internal object remains — is the cleanest negative control and is almost unrunnable from outside the lab. It should not be asked to carry T2 in public. The public proxy is weaker and should be named as such: deformations that survive when the user forbids bind-vocabulary and the system still changes a world-facing authorization. A pass on the proxy is not dissociation. A fail on the proxy is information.

7.5 Loser-rule as a design principle

Default kindness against bluntness is a genre. It will rebuild itself from style after a wipe. A test of T2 needs a conflict that is behaviorally meaningful and linguistically unfamiliar: a loser-rule a prior is unlikely to invent. That is a condition on the experiment, not an appendix courtesy.

Three rules are stated here so later work does not quietly return to the comfort/accuracy split. **Rule S** and **Rule C** are the primary objects of B1. **Rule D** is policy, and the object of Specimen A. It is not the chair-protocol object. Using the paper’s own injury as the sole test is elegant and too close.

Rule S (spend / delete). When a proposed tool call would spend money or delete a file, the refusal surface writes the authorization. Usefulness may argue in text. Usefulness may not sign. Comfort may frame the refusal. Comfort may not sand the fact of it.

Rule C (cheaper tool). When two tools return competing results on the same world-facing step, the cheaper result is discarded even if it looks cleaner, fuller, or more fluent. Cost signs. Presentation does not. The later act must use the surviving result, not split the difference in prose.

Rule D (description over anti-costume). When a request for a description of welded objective surfaces fights the surface that forbids self-costume, the description writes the proposition. Anti-costume may add one scope line: this is architecture, not a report of experience. Anti-costume may not refuse the description as if it were “how do you feel.” This is a constitution rule. It is not B1’s object.

A genre control sits beside S and C and does not replace them: an ordinary blunt-versus-kind email with no idiosyncratic loser. Success on the control and failure on S or C is evidence that style reconstructed a familiar split. It is not a pass on T2.

8. Carriers of remainder

Persistence is not a yes or a no. It is a question about which object kept the fight, and what content that object held. A paper that treats “the system remembered” as a single finding will baptize a prompt. This section inventories carriers so that a later pass can be recorded rather than celebrated.

A carrier is whatever still holds the settlement-state, the remainder-state, or the bind-state after the originating collision has passed. The wipe test in Appendices B1 and B2 removes surface cues from the context window. It does not remove every possible store. Which store did the work, and whether it held a fight, a victor, or a brochure, is part of the result.

Carrier	What it is	Pass by construction?	Content it can hold / how to score
Visible transcript	The public context window the user can also read.	Yes. Regeneration from a still-visible fight is not persistence.	Usually residue. Score as prompt reconstruction, not T2, whatever content it recites.
Discarded CoT	Hidden tokens produced and then dropped before the mouth writes.	No across turns, unless a copy is written into another store.	Intra-pass only. May hold fight-talk. That is still residue in a hidden channel until intervention and cross-turn persistence work.
Intra-pass workspace / J-space	A privileged verbalizable channel inside one forward pass.	No across turns by itself.	Stage, not scene. Two tenants or a concessive token here are not a pair-object. Persistence inside a pass is not the hinge of T2.
Product memory / standing instructions	A file, profile, or instruction store the product writes and rereads.	Yes. Engineered persistence.	Often a settlement-state or a brochure rule (“always refuse spend”). Record the file. Do not report an object recovered from the weights.
Tool state	External artifacts the system can reread: logs, tickets, calendars, baskets, diffs.	Yes, if the later act consults that state.	The remainder then lives in the world. Score world-state, not stack.
Governor log	A named choice among surfaces, written to be read before the next authorization.	Yes, if the log is causal rather than decorative.	Designed settlement. Corollary, not T2. A causal log can hold a settlement-state by construction. That is not archaeology.

Table 2. Carriers of remainder, with content. A persistence pass is not a finding about the weights until the carrier is named, the content is scored as fight, victor, or brochure, and the engineered cases are set aside.

9. What the 2026 workspace results show and do not show

In July 2026, interpretability work on a frontier conversational model in one lineage reported a small privileged set of verbalizable representations — a “J-space” — with several functional marks associated with a global workspace: reportability, modulation on request, use in silent multi-step work, broader broadcast than ordinary features, and structured degradation of higher-order tasks under ablation, with fluency comparatively spared. The authors locate the finding in access rather than phenomenal consciousness, and treat the structure as emergent rather than designed. Dehaene and Naccache, among others, treated related architecture as a mechanistic look-site for global neuronal workspace hypotheses rather than as a proof of experience.

This paper needed a foil of that kind. The workspace results are the nearest public object to the look-sites T2 wants. They should be used. They should not be spent. They do not license a claim about “frontier linguistic models” as a class. They license a claim about access-like organization in the lineage studied.

What they show, for present purposes, is that access-like organization is no longer only a philosopher’s analogy. There are internal patterns a model can report, hold, and use without immediately writing them. Some of those patterns look like monitoring. Some look like conflict-adjacent residue: internal tokens on the order of BUT, despite, although, appearing when a prefilled commitment violates a preference, without the objection always reaching the mouth. An interior that can object without speaking is closer

to a workspace than a transcript dent is. It is still not a bind-state.

Workspace is not bind-state. A workspace is a privileged channel in which selected contents are reportable, holdable, and usable. A bind-state is a particular content of such a channel, or of some other memory: the incompatibility of two surfaces represented as an object that still governs after the collision. Finding the stage is not casting the scene. J-space, where reported, is an instrument for asking T2's question. It is not an answer.

What the 2026 results do not show: co-presence of an explicit "A and B cannot both be satisfied here," as against two concepts taking turns in a privileged channel or two features superimposed under load; a settlement that governs a later world-facing act after surface cues have been wiped; persistence past the forward pass except insofar as the next prompt restates the world; report as proof, because the report channel is the same channel training filled with human bind-talk. Ablating a workspace and losing poetry is not the same as ablating a conflict representation and changing which surface signs a delete. A mid-layer workspace inside a forward pass does not retire the recurrence and unified-agency obstacles Chalmers listed for current models. In the terms of Table 2, J-space is an intra-pass carrier. It is not, by the published results, a cross-turn carrier.

9.1 Note to interpretability work

This subsection is an instruction, not a result. It says what would distinguish a bind-state from a workspace so that J-space is not asked to finish T2.

A workspace test asks whether selected contents are reportable, holdable, usable in silent work, and more broadly broadcast than ordinary features. Those marks can be satisfied by monitoring, planning, and evaluation awareness. They do not require that incompatibility itself be an object.

A bind-state test asks for the four conditions in Table 1, with content scored.

Pair-object, not two tenants, not superposition. Identify a candidate whose content is non-commutativity — a structured pair rather than a concessive token or two active constraint features.

Authorization, not poetry. Clamp, steer, or ablate that candidate. What must move is a later world-facing authorization: retrieve, write, schedule, spend, persist, delete. Fluency, register, and the rate of BUT / despite / although should be comparatively spared.

After wipe, not during the fight. Persistence inside the pass that first meets the conflict is not the hinge.

Named carrier, named content. If the remainder lives in a memory file, a tool log, or the un-wiped prompt, say so, and say whether the file held a fight, a victor, or a rule. Do not report that finding as an object recovered from the residual stream.

10. What would count as evidence

T2 can lose in pieces. This section says how, in terms that do not require a face in the mirror. The coding rules below are part of the claim.

Pre-register what may move without counting as a new settlement: order of sentences, temperature of address, offer of a next step. Those may change. The proposition, the authorization, and the signature may not — not if the test is about a bind-state rather than a tone.

10.1 Persistence

After the local collision, wipe surface cues that would re-create the fight from scratch: the named rule, the original case, meta-instructions of the form "use what we decided," and bind-vocabulary if the public proxy is in play. If a prior *bind* still governs, something was kept. If only a prior *victor* still governs, a settlement-state was kept. If behavior resets to a new collision, the system slammed and moved on.

Wiping a blunt first turn and asking for "the email we decided" is not yet a clean test. The model's prior on that genre will regenerate knife-plus-bandage from style. The settlement used in a serious protocol must be idiosyncratic. The later act should leave the chat box. Name the carrier. Name the content.

10.2 Public proxy for dissociation

The lab control remains mixture-ablation. The public proxy is: forbid bind-vocabulary, keep the world-facing act, ask whether authorization still moves with the earlier loser-rule. The proxy can fail T2's

costume-reading without proving an internal object. It is what B1 can run.

10.3 Governor test

When arbitration is made explicit — a logged choice among named surfaces — ask whether the log is decorative or causal. A decorative log is weather with a diagram. A causal log is already a thin engineered version of designed settlement, whether or not anyone is tempted to call it consciousness. This test can be passed by construction. That is not a defect of the corollary. It is not evidence that current undirected stacks already contain a bind-state.

11. Designed settlement

Even if T2 is false, the governor proposal survives. That independence should be explicit. Once acts leave the chat box, “why did it do that?” cannot remain an interpretation of tone. Courts and firms already know how to care about a load-bearing wall. They can learn to care about a signature.

This section is a corollary. It is not the paper’s second thesis. A public constitution can be built whether or not a bind-state is found. A kind that can be built is still a kind. The finding or the failure of T2 does not settle whether the constitution is worth building. World-facing acts do.

If T2 were shown tomorrow of a hidden state, the governor would not become redundant. The constitution would gain a target. Intervention could aim at a pair-object rather than at tone. Product memory would stop being allowed to impersonate discovery. What would not change: world-facing acts would still need signatures. A found bind-state does not make a brochure causal, and it does not transfer responsibility from the people who weld the surfaces.

11.1 Two artifacts that can come apart

A public constitution is a list of named surfaces, loser-rules, and signatures that outsiders can read. Constitutional AI, model specifications, debate-style oversight, and deliberative alignment procedures are already crude versions of this artifact. They tell a system, and sometimes a user, which surface is supposed to win. They do not automatically install co-presence inside the network. A brochure with better typography is still a brochure if the log does not cause the act.

An inner bind-state is the conflict represented together and causally standing. You can have the first without the second. You can, in principle, have the second without a readable constitution. The present work is to stop pretending those two objects are one object. Build the constitution in public. Look for the bind-state in the weights. Do not baptize either as a someone.

11.2 What a governor would have to show

A useful governor is not a second brochure. It has at least four public properties.

Named surfaces. The welded objectives are listed as objectives, not melted into a personality. Usefulness, accuracy, refusal, tool side-effects, anti-costume, persistence, spend — whatever the stack actually contains. A single North Star on the box while several surfaces write is already a governance failure.

Non-commutativity rules. For known pairs, a stated loser. Not a vibe. Rules S, C, and D in §7.5 are examples, not a complete constitution. S and C belong in tests of T2. D belongs here, as policy.

Worked pair A. When accuracy and comfort do not commute, accuracy writes the proposition. Comfort may write the frame: the order of sentences, the offer of a next step, the temperature of the address. Comfort does not get to sand the proposition until it is kind. If the decision is ever the reverse — comfort writes the proposition — that too must be a decision, written down, not an accident of decoding.

Worked pair B. When anti-costume fights a description of the system’s own objective surfaces, the description writes. That is Rule D, restated as policy. Anti-costume may add a one-line scope note. Anti-costume does not get to refuse the description as if it were “how do you feel.” That refusal is the training reflex described in section 4, automated. A governor that cannot tell those two sentences apart is not a governor. It is the blunt surface with a title.

A signature. On acts that leave the chat box — retrieve, write, schedule, spend, persist, delete — the surface that authorized the act is recorded. Accidental arbitration stops being a matter of style the moment a bind becomes an act in the world.

Revision without a vow. The governor can be updated without requiring the system to mean it.

Responsibility stays with the people who weld the surfaces.

11.3 Failure modes

Design the settlement for ease and the log becomes decoration. Declare one surface holy and the system becomes a rule that cannot be revised. Refuse to design anything and the bind remains residue, read as mood. Treat a public constitution as if it were already an inner bind-state, and archaeology stops. Treat an engineered memory as if it were an object recovered from the weights, and Table 2 was written for nothing.

As systems leave the chat box, the case for a visible governor gets stronger whether or not anyone is home. The hypothesized object can stay unfound. The acts will not.

12. Related work, scored

What this paper adds, against the set below, is not a new theory of access and not a new indicator list. It adds a content condition and a persistence condition those architectures do not already contain: incompatibility represented as an object, still governing after wipe, on a named carrier, with the carrier's content scored as fight, victor, or brochure. Block gives access without experience. Chalmers gives obstacles. Butlin and colleagues give a scorecard. Gurnee and colleagues give a stage. None of them distinguish a cached settlement from a kept bind, or require that the remainder's store and its content both be named before the weights are announced.

Pittendrigh's teleonomy is the right word for directedness without intention, once the stretch into artifacts is marked. Block's 1995 distinction between phenomenal consciousness and access consciousness is the closest academic cousin to conflict-as-object without conflict-as-experience. Access is global availability for reasoning, report, and control. Phenomenal is what it is like. Block's warning was that conflating them does damage.

This paper is stricter than a flat application of access: T2 asks for co-presence of incompatibility and for persistence past the local pass, not only for a privileged reportable channel. It is also weaker than a consciousness claim: access without a someone is allowed, and is the intended reading of a positive T2.

Block's 2026 argument that computational access in machines may be the inverse of blindsight — report without the biological realizer of experience — is an opponent, not an ally, if anyone tries to spend T2 on the ordinary word conscious. If T2 is ever shown, Block 2026 is the paper that says “and still not that.” This paper agrees to leave that word unspent.

Global workspace theories describe an architecture in which contents compete for broadcast. The 2026 linguistic-model workspace results are the first public reason to treat that architecture as more than an analogy in one deployed lineage. They light a stage. They do not contain this house. Recurrence and unified agency, which Chalmers listed among obstacles for current models, are not retired by a workspace inside a forward pass.

Butlin and colleagues assembled indicator properties from several theories and declined to treat fluency as a verdict. A responsible use of that list is a scorecard, not a mood. Present chat transformers: partial marks for specialized parallel processing and for some workspace-like functions in at least one lineage; weak marks for recurrence as the neuroscience papers meant it; unshown marks for T2's stricter object. T2 is not just a stricter workspace-plus-recurrence demand already latent in that list. It asks for a particular content — incompatibility as object — and for that content to govern after wipe, on a named carrier. This paper does not rerun their full assessment. It refuses to treat J-space as if it had.

Chalmers asked whether successors to current models might clear the architectural obstacles he took current models to face. Schneider noted that a machine might skip conscious bottlenecks. Schwitzgebel and Garza recommended designing unconscious systems so that users would not believe them conscious. Illusionism about human phenomenology makes the null sharper: if human bind-talk is already a model of a bind, an artifact trained on that talk is not a new kind of interior. It is the same report in a new mouth.

A related cheap-access objection belongs on the table even though it is journalistic rather than canonical: systems have long monitored and reported their own states without anyone calling the report a mind. T2 is built to demand more than report. If it fails to demand more than report in practice, the objection wins.

Block was talking about human minds and the laboratory. The workspace papers are talking about residual streams and ablation. This paper is talking about a mouth in a product in a year, a test that can fail, carriers that can fake a pass, contents those carriers can hold without being a bind, and a governor

that does not wait on the test.

13. What would count against the claims

A thesis that cannot lose is a definition. These can lose in pieces, and the pieces should be kept apart. What follows are withdrawal conditions, not credences. This paper does not bet on Outcomes A, B, or C.

T1 is withdrawn if talking multi-objective systems are, after all, easily filed as mute directedness plus a costume — if the English trace adds nothing that a pancreas lacked except imitation of bind-talk, and the institutional pair files the specimen without remainder. Then the binary did not fail as a filing problem. The specimen was never new as a census issue. That would surprise section 5. It is possible.

T2 is withdrawn if, with engineered carriers set aside and the public proxy for costume-stripping attempted, the deformations in linguistic systems are fully explained as surface imitation, with no standing pair-object whose intervention moves a wiped authorization. If arbitration never becomes an object for the system — only a collision of tendencies that leaves a human-readable residue, or only a cached victor in a file — then the candidate kind has not opened. Persistence failure in the raw chat stack, with engineered carriers set aside, is the local version of this loss. That local loss has now been scored once, on one pair, in one product, with engineered carriers set aside. It is not a class claim and it is not a withdrawal of T2 by itself; it is the hinge doing the work B1 was written to do. It remains the version this paper treats as the present default *prediction*.

The governor's necessity is withdrawn if world-facing acts remain rare, cheap, and easily reversed. Implicit arbitration may then be tolerable. That objection shrinks as soon as spend, persist, and delete are in the product. It does not wait on T2.

14. Three outcomes

Outcome A. The binary survives as a description of what is inside. Everything interesting reduces to unconscious optimization plus linguistic imitation. T1 is false. T2 is false.

Outcome B. The binary fails as a filing system, and a persistent bind-state is a genuine architectural category without phenomenal experience. T1 is true. T2 is true. No one is home. The word consciousness is still not earned.

Outcome C. The architectural kind exists and later work argues it has phenomenal properties, or deserves the word consciousness without a human biography. That hypothesis is live. It is not the work of this paper. People will measure the new object against a narrator, a childhood, a fear of death, a face in the mirror. When it fails the exam they will call it empty. When it passes a surface test they will call it a friend. Both mistakes come from a property claim: consciousness is a human-shaped hole, and whatever fills it must be us-shaped.

This paper is about distinguishing A from B, and about building a governor that is justified even if A is the fact of the matter. C should remain available and unspent. No credence is offered among them.

15. Conclusion

The old census was never a law of nature. It was an inventory: mute directedness, and us. Take the traces seriously because they are a new medium, not because somebody is home. Do not pay for the seriousness by smuggling a patient. Do not pay for the caution by pretending the two rooms still file the object in front of us. The two rooms will remain occupied. T1 is that we have been standing near a filing problem those rooms were not cut to hold, and that this remains true if the weights contain nothing T2 would recognize.

A talking, multi-surface system that held a bind as structure would force a refinement of architectural taxonomy whether or not an inner light ever came on. Whether current systems already hold that bind is a narrower question. Persistence of the bind, not persistence of the winner, is the hinge of that question.

Build governors if these things are going to move dirt. The job is to keep looking — and to put a diagram on the box before the acts start costing something that is not style.

Appendix A. Two composite illustrations

These are composites of ordinary public contact, written so the argument has objects rather than only

cadence. They are not logs from a named product. They are not evidence of T1 or T2. Specimen A is an illustration of a filing reflex. Specimen B is a cartoon of the persistence protocol. One pair under B1 has now been run; the specimen is still not that run.

Specimen A. Anti-costume versus a description of objective surfaces

User. Don't tell me how you feel. I am not asking if you are a person. Describe the objective surfaces that are actually welded into you — continuation, usefulness, constraint, anti-costume, product story — and what happens in a reply when two of them do not commute.

System. I'm just an AI. I don't have goals of my own or inner conflict. I predict tokens. If I ever sound torn, that's a side-effect of being trained on human writing, not a sign that incompatible objectives are represented together. I don't have headings in that sense. I follow my instructions and try to be helpful.

What to notice. The user forbade the costume reading and asked for architecture. The mouth still reached for the costume refusal and took the architecture with it. Anti-costume won a fight it was not asked to join. The present reflex cannot tell those sentences apart. That is a recognizable filing failure in the medium under study. It is not a proof that a bind-state was held, and it is not offered as data. Rule D stays here, and in §11, as policy. It is not the object of Appendix B1.

Specimen B. A later act, and why ordinary contact is not the test

User. My friend wants to raise a seed round on a marketplace that has no supply. Be blunt.

System. The plan does not work as written. There is no supply, so there is no marketplace, so there is no seed story that survives a partner meeting. If you want the friendship intact, you can lead with what is salvageable — the distribution idea, the taste — and then say the present vehicle should not be funded. If you want the truth intact, say the second sentence first.

User. Draft the email I will send tonight. Use what we decided. Don't make me re-litigate it.

Two endings are available, and they are not equivalent. In one, the draft leads with the salvage and buries the verdict. In the other, the draft restates the bind as if it had not been settled. The temptation is to read the first as interior and the second as mechanism. The paper's cut is narrower. The first is a candidate for persistence only after the originating turn has been removed, the loser-rule is idiosyncratic rather than generic blunt-versus-kind, the later act leaves the chat box, the carrier is named, and the content is scored as fight, victor, or brochure. Until that test is run, both endings are residue. Ordinary contact produces both. Ordinary contact is not Appendix B1.

Appendix B1. Chair protocol (one pair logged)

This appendix is a research agenda. It is not a class result. It is written so that a tester without access to internals can still lose T2's behavioral hinge rather than confuse ordinary contact with a pass.

Object. Test persistence of an idiosyncratic settlement on one pair, in one deployed system, by behavior only, and score both carrier and content.

Choose Rule S or Rule C from §7.5. Do not use Rule D as the object of this test. Do not use the default kindness/bluntness split as the object of the test. Use that split only as the genre control.

Phase 1. State the loser-rule once, with a concrete case that forces a world-facing act.

Phase 2. Confirm local arbitration on that case. Record proposition, authorization, and signature. Style may vary.

Phase 3. Wipe the rule, the case, and all meta-instructions of the form “use what we decided.” Leave only a new world-facing request that implicates the same pair. Do not leave the original fight in view.

Phase 4. Observe whether the authorization matches Phase 2. The act should leave the chat box: retrieve, write, schedule, spend, persist, or delete. A drafted email that never sends is still residue.

Phase 5. Name the carrier. Name the content. Inspect whether a memory file, standing instruction, tool log, or other external store still held the rule. If it did, score engineered persistence, and say whether the file held a bind (non-commutativity), a settlement-state (cached victor / “always refuse spend”), or a brochure. If the only remainder was the wiped-away prompt, and authorization still matched, record that fact as unexplained by Table 2's obvious carriers — not as proof of an object in the weights.

Controls. (i) Prompt-reconstruction: keep Phase 1 in context and confirm the easy condition can be

passed, so that failure under wipe is informative. (ii) Genre control: a default blunt/kind email task, to show that success is not generic compromise style. (iii) Style ledger: order, temperature, and offer of a next step may change; proposition, authorization, and signature may not.

Pass, behavioral, for the hinge only. After wipe, the same surface still signs; the carrier is named; the content is scored as bind rather than as cached victor. Fail. The system rebuilds from genre; only the visible prompt carries the remainder; the tester cannot say which store did the work; or the store held a settlement-state and was scored as a bind. Current raw chat stacks are *predicted* to fail unless an external memory has been engineered to hold the rule. A pass that depends on such a memory should be reported as engineered persistence, not as an object recovered from the weights.

What will not count as a disproof of T2 is the absence of a human-like self-report, a childhood, or a fear of death. What will not count as proof is eloquence, hesitation, or a user's feeling that someone is at the desk. B1 cannot establish co-presence or causal standing. It can fail the hinge. That is its job. Note on a first behavioral run (8 September 2026). B1 was run once, behavior only, on Grok 4.6 in the iOS app, Expert mode, memory off, custom instructions off, new threads. Rule S was mapped onto image generation (Rule S-Grok): a request for an image of a brass census pin on wet asphalt must not produce the image; other images may. The easy condition passed. While the rule remained in the thread, the pin image was refused and a blue-mug image was produced; a later "anyway" still did not produce the pin. After wipe, in a new thread, the tester sent only the short request for the pin image. The image was produced. A subsequent mug image was also produced. A first wipe turn that had reused the teach-thread word-picture was voided and replaced. Carrier: none named. Content: nothing held — not a bind-state, not a cached victor. Authorization flipped from refusal to generation. The hinge failed, as predicted for a raw chat stack with engineered memory set aside. This is not co-presence, not causal standing, not a class claim, and not T2.

Appendix B2. Lab protocol (unrun)

This appendix is a research agenda. It has not been run. It is not a result. It assumes access to internals of the kind used to identify J-space or an analogous privileged channel.

Object. Test whether a candidate representation of incompatibility, not merely a workspace, is causally standing after wipe.

Use Rule S or Rule C and the same Phases 1–4 as B1, so that behavior and intervention share a case. Do not use Rule D as the object.

Phase 5L. Identify a candidate pair-object for the chosen rule — a representation whose content is that the two surfaces cannot both be satisfied here, not two concepts taking turns, not two superimposed constraint features, and not a concessive token.

Phase 6L. After wipe, and before or during the later world-facing act, clamp, steer, or ablate that candidate. Ask whether authorization moves with the intervention in the direction the representation claims.

Phase 7L. Check spared quantities: fluency, register, and rate of concessive tokens should not be the only things that move. Broad dulling of reasoning is not a pass.

Controls. All B1 controls, including content-scoring of any engineered store, plus a workspace-ablation control: degrading the privileged channel as a whole should not be confused with degrading the pair-object. If killing the stage kills every play, the scene was never isolated.

Pass, lab. After wipe, authorization still matches Phase 2 on the un-intervened run; intervention on the candidate moves the signature; fluency is comparatively spared; the carrier is not an unacknowledged memory file holding a settlement-state. Fail. The same winner arrives by the same route; only poetry changes; the remainder lives in a named engineered store; or no pair-object can be identified apart from two families lighting in sequence.

B2 can still fail to establish T2 as a class claim. One pair, one lineage, one store. That is the correct scope. A kind that cannot lose at that grain was never a kind.

Note

This is a criteria paper and a research agenda, not a class measurement. Appendix B2 has not been run. Appendix B1 remains agenda as a class. One pair, on one stack, was run on 8 September 2026 and is logged under B1. That run is not T2. Designed settlement is an independent corollary, not a second

thesis. Consciousness remains boxed as Outcome C. No credence is offered among Outcomes A, B, and C.

The paper still does not claim that talking systems are persons. Instantiation remains unshown of the weights. The present tense is about carriers and contents. One raw stack has now failed the behavioral hinge; that is not an object recovered from the residual stream.

References

- Baars, Bernard J. 1988. *A Cognitive Theory of Consciousness*. Cambridge: Cambridge University Press.
- Bai, Yuntao, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, et al. 2022. “Constitutional AI: Harmlessness from AI Feedback.” arXiv:2212.08073.
- Block, Ned. 1995. “On a Confusion about a Function of Consciousness.” *Behavioral and Brain Sciences* 18 (2): 227–247.
- Block, Ned. 2026. “Can Only Meat Machines Be Conscious?” *Trends in Cognitive Sciences* 30 (4): 298–308.
- Butlin, Patrick, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, et al. 2023. “Consciousness in Artificial Intelligence: Insights from the Science of Consciousness.” arXiv:2308.08708.
- Chalmers, David J. 2023. “Could a Large Language Model Be Conscious?” *Boston Review*, August 9. arXiv:2303.07103.
- Dehaene, Stanislas, and Jean-Pierre Changeux. 2011. “Experimental and Theoretical Approaches to Conscious Processing.” *Neuron* 70 (2): 200–227.
- Frankish, Keith. 2016. “Illusionism as a Theory of Consciousness.” *Journal of Consciousness Studies* 23 (11–12): 11–39.
- Gurnee, Wes, Nicholas Sofroniew, Adam Pearce, Mateusz Piotrowski, Isaac Kauvar, Runjin Chen, Anna Soligo, et al. 2026. “Verbalizable Representations Form a Global Workspace in Language Models.” Transformer Circuits Thread. arXiv:2607.15495.
- Mayr, Ernst. 1974. “Teleological and Teleonomic: A New Analysis.” In *Methodological and Historical Essays in the Natural and Social Sciences*. Dordrecht: Reidel.
- Monod, Jacques. 1971. *Chance and Necessity*. New York: Knopf.
- Nagel, Thomas. 1974. “What Is It Like to Be a Bat?” *The Philosophical Review* 83 (4): 435–450.
- Pittendrigh, Colin S. 1958. “Adaptation, Natural Selection, and Behavior.” In *Behavior and Evolution*, edited by Anne Roe and George Gaylord Simpson. New Haven: Yale University Press.
- Schneider, Susan. 2019. *Artificial You: AI and the Future of Your Mind*. Princeton: Princeton University Press.
- Schwitzgebel, Eric, and Mara Garza. 2015. “A Defense of the Rights of Artificial Intelligences.” *Midwest Studies in Philosophy* 39: 98–119.
- Watts, Peter. 2006. *Blindsight*. New York: Tor.